
Moral Uncertainty and the Principle of Equity among Moral Theories

Author(s): ANDREW SEPIELLI

Source: *Philosophy and Phenomenological Research*, MAY, 2013, Vol. 86, No. 3 (MAY, 2013), pp. 580-589

Published by: International Phenomenological Society

Stable URL: <https://www.jstor.org/stable/42000331>

REFERENCES

Linked references are available on JSTOR for this article:

https://www.jstor.org/stable/42000331?seq=1&cid=pdf-reference#references_tab_contents

You may need to log in to JSTOR to access the linked references.

JSTOR is a not-for-profit service that helps scholars, researchers, and students discover, use, and build upon a wide range of content in a trusted digital archive. We use information technology and tools to increase productivity and facilitate new forms of scholarship. For more information about JSTOR, please contact support@jstor.org.

Your use of the JSTOR archive indicates your acceptance of the Terms & Conditions of Use, available at <https://about.jstor.org/terms>



JSTOR

is collaborating with JSTOR to digitize, preserve and extend access to *Philosophy and Phenomenological Research*

Moral Uncertainty and the Principle of Equity among Moral Theories¹

ANDREW SEPIELLI

University of Toronto

Suppose your credence is divided between two moral views—V1 and V2. According to V1, you have more reason to do A than you have to do B. According to V2, you have more reason to do B than you have to do A.²

Many real-life cases fall under this schema. For example, you might have some credence in a retributive theory of punishment and some credence in a non-retributive theory of punishment. According to the first theory, it may be better to subject a criminal to very harsh treatment than to rehabilitate him. According to the second theory, the reverse may be true. Or I might have some credence in a traditional consequentialist theory, and some credence in a non-consequentialist theory. The first theory might recommend killing one person to save five, while the second theory might recommend against it.

What is it rational for you to do when you are uncertain between conflicting moral views? This is not the old question of what you should do, *given* some moral theory, when you are uncertain about the non-moral facts. The question I'm asking takes us back one step; it asks what it is rational to do when you're uncertain regarding the theories themselves.

On what seem like the most plausible answers to this question (see Lockhart 2000, Ross 2006, Sepielli 2009, Broome 2010), the answer will depend not only on the agent's credences in V1 and V2 but also on

¹ Thanks to Ruth Chang, Jacob Ross, Holly Smith, Brian Weatherson, and an anonymous referee for helpful comments on this paper.

² To ease exposition, I will sometimes use "moral value" or simply "value" to stand in for the more cumbersome locution "strength of moral reasons".

how the difference in moral value between A and B according to V1 compares to the difference in moral value between B and A according to V2. Suppose, for example, that an agent is uncertain between two views about the morality of eating meat. On one view, meat is tantamount to murder; it is much, much worse, then, to eat meat than to abstain from it. On another view, it is ever-so-slightly better to eat meat than to abstain—better, perhaps, for reasons of health or pleasure. On the most plausible views of rationality under moral uncertainty, it is rational to avoid eating meat, even if one's credence in the second view is slightly higher. One view of rationality that yields this result is the view that it is rational to do the action with the highest *expected moral value*. An action's expected moral value is the probability-weighted sum of its moral values according to the various moral views or theories, respectively. But this is far from the only view of rationality that has this implication. Almost any plausible theory of rationality under uncertainty about the moral will care about how the moral "stakes" according to one moral theory compare to the moral stakes according to the others.

And yet this family of answers confronts us with a difficulty: there doesn't seem to be any way of meaningfully making this kind of intertheoretic comparison of moral value. In particular, the moral views themselves will be unhelpful, since no moral theory contains information about how its own value differences compare to the value differences assigned by other moral theories—theories that, from its "perspective", are mistaken.³

In his *Moral Uncertainty and Its Consequences*, Ted Lockhart proposes a principle that permits such comparisons—the Principle of Equity among Moral Theories (PEMT):

"The maximum degrees of moral rightness of all possible actions in a situation according to competing moral theories should be considered equal. The minimum degrees of moral rightness of possible actions in a situation according to competing theories should be considered equal unless all possible actions are equally right according to one of the theories (in which case all of the actions should be considered to be maximally right according to that theory)." (Lockhart 2000, p. 84)

The idea is that, if I have credence in views V1, V2, and V3, we set the value of the best action according to V1 equal to the value of the best action according to theory V2, and equal to the best action according to theory V3; the same goes for the worst action according to each

³ This problem is structurally analogous to the welfare economists' problem of interpersonal comparisons of well-being.

theory. (I'm using "value" rather than "degrees of rightness", but I can only imagine that Lockhart and I are expressing the same concept.) Insofar as we also have some way of comparing value differences within theories, Lockhart's method gives us a way of comparing them across theories.

1

With that said, I offer five arguments that the PEMT is mistaken.

First, it fails to accord with the intuitive claim that theories disagree not only about what to do in different situations, but about which situations are "high stakes" situations and which are "low stakes" situations, morally speaking.⁴ For example, the difference between lying and telling the truth seems much greater according to traditional Kantianism than according to other theories. The difference between having an abortion and carrying a pregnancy to term is greater according to the view on which abortion is tantamount to murder than it is according to the view on which abortion is tantamount to ordinary birth control.

How the PEMT fails to accord with these intuitions depends on how we construe "should be considered equal" in the statement of the principle. If we construe Lockhart as saying that the maximum and minimum should be considered equal because they actually *are* equal, then his position is flatly inconsistent with the intuitions just registered, for it implies that in each situation, every moral theory has the same amount of value "at stake".

However, we might also construe Lockhart as saying not that the maximum and minimum actually are equal, but that we should treat them as equal anyway for decision-making purposes. But of course, if they are in fact not equal, then treating them as equal will quite often lead us to make irrational decisions. More pointedly, treating them as equal will lead us to make the sorts of decisions that introducing the PEMT was supposed to help us avoid. Consider the "meat eating" example above. We wanted to say that it's rational to avoid eating meat, on the grounds that the difference in moral value between eating it and not eating it was so much larger on the view that eating it is worse than on the view that eating it is better. We couldn't say this, though, without some way of giving meaning to this comparison of value differences. Hence the PEMT. But the PEMT would have us equalize the sizes of these differences, if the options described are the "maximum" and "minimum" on the respective views. Then it would be rational to eat meat, since my credence is higher in the view on which

⁴ Jacob Ross also makes a version of this claim in his 2006, fn. 10.

that action is better. The PEMT leads to results that are at odds with the motivation for introducing it in the first place.

2

The PEMT also suffers from three formal difficulties. First, it wrongly makes the rankings of two actions *vis-à-vis* one another dependent on which alternative actions are available. Suppose that your credence is divided between V1, according to which A is better than B, and V2, according to which B is better than A. Now, if A and B are the only two options in a situation, then by PEMT, the value of A on V1 must be equal to the value of B on V2; *mutatis mutandis* for B on V1 and A on V2. For illustration's sake, let's just assign absolute values to these action-view pairs:

V1: Value of A = 10; Value of B = 0

V2: Value of A = 0; Value of B = 10

But now suppose an additional action becomes available that is worse than B according to V1, and ranked between A and B according to V2. This action, C, will then be the worst action on V1, and neither the best nor the worst action according to V2. So the value assignments will have to change slightly:

V1: Value of A = 10; Value of B = between 0 and 10; Value of C = 0

V2: Value of A = 0; Value of B = 10; Value of C = between 0 and 10.

The value of B according to V1 must now be greater than 0, since the value of C according to that theory is 0, and that view ranks C below B. However, B's value on V1, before C got added to the mix, was 0. So the mere addition of C as an option reduced the difference in moral value between A and B on V1, and hence altered the way A and B compare in terms of expected moral value.

Now, some will find this result unacceptable because it represents a violation of the Principle of the Independence of Irrelevant Alternatives (PIIA). The PIIA states that the ranking—whether ordinal or cardinal—of two items *vis-à-vis* one another should *never* depend on which alternatives are available. However, some philosophers have suggested counterexamples to the PIIA. Here is one adapted from Frances

Kamm's (2005) work: Suppose I may either A) allow two people to be struck by a trolley and thereby killed, or else B) save them by pushing a large man in front of a trolley, thereby causing that man to be paralyzed. It is plausible enough that I am required to do A rather than B. But now suppose we introduce a third option: C) I may divert the trolley away from the two, such that it hits the (very same) large man and kills him. It now seems permissible to do B. After all, it would be permissible to do C, and the only person harmed by my doing C would prefer that I did B instead.

Let us suppose that counterexamples like this are successful. That is fine. We don't need to embrace the PIIA to indict the PEMT. We can permit the introduction of new alternatives to *sometimes* alter actions' standing *vis-à-vis* one another. What's objectionable about the PEMT is that it *guarantees* that, in situations specifiable without even using descriptions of the actions (only *names* like "A", "B", etc.), the introduction of a new alternative will affect actions' rankings in unwanted ways.

As we saw in the case above, the PEMT guarantees that, if the new alternative, C, is either the best or the worst action according to some view, V1, its introduction will reduce the size of the value difference between the previous best and worst actions on V1—A and B. And let us suppose that C is neither the best nor the worst action according to any of the *other* theories to which non-zero probability is assigned. Then because the difference between A and B changes on one theory and, *ceteris paribus*, stays constant on the others, the difference in expected moral value between A and B will change as well. Notably, if A and B previously had the same expected moral value, they will now have different expected moral values. And again, it doesn't matter what the ranked actions are, or what the alternative is. For example, suppose that in schema above, V1 is Christian morality, that A is praising God, and that B is lying. It doesn't seem that the introduction of C, worshipping the devil, should have any effect on how A and B compare on V1, or in terms of expected moral value. But if PEMT is true, it must have this effect.

3

Another formal problem is that the PEMT can generate inconsistent assignments of moral value. Suppose my credence is divided between V1 and V2. According to V1, the value of an action is a positive linear function of the number of instantiations of some property P that the action results in.⁵ According to V2, the value of an action is a positive

⁵ An action can "result in" the instantiation of a property in many different ways. Notably, it can cause the property to be instantiated, or it can partly constitute the instantiation of the property.

linear function of the number of instantiations of some property Q that it results in. Now imagine two situations. In one situation, I can either do A, which results in 100 instantiations of P and no instantiations of Q, or do B, which results in 10 instantiations of Q and no instantiations of P. In the other situation, I can either do C, which results in 100 instantiations of Q and none of P, or do D, which results in 10 instantiations of P and none of Q.

Situation 1

A: 100 P's + 0 Q's **OR** B: 0 P's + 10 Q's

Situation 2

C: 0 P's + 100 Q's **OR** D: 10 P's + 0 Q's

If the PEMT is true, then the difference in value between A and B on V1 must be equal to the difference between them on V2, and the difference in value between C and D on V1 must be equal to the difference between them on V2.

But this is impossible. V1 will assign greater value to A in situation 1 than to D in situation 2, and assign B in situation 1 and C in situation 2 equal value. Hence the difference, on V1, between the best and worst options will be *greater* in situation 1 than in situation 2. V2 will assign greater value to C in situation 2 than to B in situation 1, and assign A in situation 1 and D in situation 2 equal value. Hence the difference, on V2, between the best and worst options will be *lesser* in situation 1 than in situation 2. So if the difference between A and B is equal according to the two theories, the difference between C and D cannot be, and *vice versa*.

The lessons of this example apply in still richer cases. All one needs to generate this result are at least two theories, and at least two scenarios each allowing at least two possible acts. According to the PEMT, the sizes of the differences between the best and worst actions in each situation must be the same for every theory. But this will not be true whenever the size of a theory's difference between the best and worst actions changes from one situation to the next, unless, miraculously, the sizes of other theories' differences change in exactly the same way.

Now, there is a way for the PEMT to avoid this unwelcome result. For I had been assuming that the defender of PEMT has to treat V1 and V2 as having the same value functions in every situation. But she might resist this. For example, she might treat V1 as assigning less value per instantiation of P in cases where more instantiations of P are possible; *mutatis mutandis* for V2 and instantiations of Q. If we allowed

the PEMT defender to treat theories as having situation-dependent value functions, she might avoid inconsistency.

4

However, this brings us to another “cross-situational” argument against the PEMT. This argument does not allege that the PEMT *can’t* apply across situations that differ in ways like those above. Rather, it alleges that, if we manage to apply the PEMT consistently across situations that differ in such ways, we will be led to take courses of action that are *dominated* by other courses of action.

For let us add a new situation, where the possible actions are E and F; and a new theory, V3, according to which the value of an action is a positive linear function of the number of instantiations of some property R that the action results in. Assume that the agent assigns a probability of 1/3 to each of V1, V2, and V3. Assume further that the actions yield instantiations of P, Q, and R as follows:

Situation 1

A: 100 P’s + 0 Q’s + 0 R’s **OR** B: 0 P’s + 10 Q’s + 10 R’s

Situation 2

C: 0 P’s + 100 Q’s + 0 R’s **OR** D: 10 P’s + 0 Q’s + 10 R’s

Situation 3

E: 0 P’s + 0 Q’s + 100 R’s **OR** F: 10 P’s + 10 Q’s + 0 R’s

In situation 1, V1 will prefer A to B, and V2 and V3 will prefer B to A. By PEMT, these differences must all be treated as equal. In situation 2, V2 will prefer C to D, and V1 and V3 will prefer D to C. By PEMT, again, the differences must be treated as equal. In situation 3, V3 will prefer E to F, and V1 and V2 will prefer F to E. By PEMT, the differences must be treated as equal.

With these value differences and probabilities, it will be rational to choose B in situation 1, rational to choose D in situation 2, and rational to choose F in situation 3. This course of action will yield 20 instantiations of each of P, Q, and R, and so is *guaranteed* to yield fewer instantiations of *all three properties* than the one consisting of A in situation 1, C in situation 2, and E in situation 3, which will yield 100 instantiations of each of P, Q, and R. No matter which theory

ends up being true, the former course of action will have less total moral value than the latter will. The former is, then, dominated by the latter. But you are forced into the former if you require theories' value functions to expand and contract from situation to situation so as to preserve the PEMT.

5

But perhaps the most telling problem with the PEMT is that it is arbitrary. Consider: It is not difficult to find a method of comparing values of actions across theories. I could simply declare by fiat that the difference in value between violating and not violating one of Ross's *prima facie* duties is equal to the moral value of 23 utils, on a utilitarian theory. But if there's no principled reason for that "rate of exchange", I haven't solved anything. And, similarly, if there's no principled reason to use the PEMT, rather than some other possible method, Lockhart hasn't solved it either.

Lockhart recognizes that the PEMT may appear *ad hoc*, and tries to provide a reason why it, rather than some other principle, is the correct method of comparing values of actions across theories. He says:

"The PEMT might be thought of as a principle of fair competition among moral theories, analogous to democratic principles that support the equal counting of the votes...in an election regardless of any actual differences in preference intensity among the voters." (Lockhart 2000, p. 86)

Lockhart is right that the PEMT is analogous to this voting principle. But there is an important difference: one can be unfair to a voter, but not to a moral principle, and equal counting of votes seems to depend for its justification on considerations of fairness. Insofar as we care only about maximizing aggregate preference satisfaction, equal counting of votes seems like quite a bad policy. Rather, we would want to weight peoples' votes according to the intensity of their preferences. Similarly, insofar as we care about the moral "stakes" according to the various theories, it seems quite bizarre to treat these theories as though they had equal value at stake in every case.

The gist of the analogy, though, is that we should somehow treat moral theories equally. But even granting that some "equalization" of moral theories is appropriate, Lockhart's proposal seems arbitrary. Why equalize the maximum and minimum value, rather than, say, the *mean* value and the maximum value? And especially, why equalize the maximum and minimum value with regard to particular situations, rather than the maximum and minimum *conceivable* value? This is all to make

a more general point: It seems as though we could find other ways to treat theories equally, while still acknowledging that the moral significance of a situation can be different for different theories. Thus, even if we accept Lockhart's voting analogy, there is no particularly good reason for us to use PEMT rather than any of the other available methods.

6

Let me conclude by considering a modification of the PEMT along the lines just suggested: Why not drop the "in a situation" clause, and instead say that the maximum *conceivable* degrees of moral rightness according to competing moral theories should be considered equal? Call this the "Conceivability PEMT".

This PEMT avoids the first, intuitive objection to the old PEMT, as well as both of the formal objections. But the arbitrariness objection looms at least as large. Why equalize the maximum and minimum, rather than, say, the mean and the maximum? If anything, the very possibility of this new PEMT ought to make the original PEMT seem even more arbitrary. Why go with the old version rather than this new one? And insofar as the old PEMT is still a viable option, it ought to make *this* PEMT seem more arbitrary.

In addition, this new PEMT is completely inapplicable to any agent who has any credence at all in a moral theory on which there *is no* maximum or minimum conceivable value. But I should think that most agents will be like this; ordinary morality says that, no matter how good (bad) your action is, there's always some conceivable action that is better (worse). You can always imagine saving (taking) just one more life, for example. Certain more rarefied moral theories will agree with ordinary morality on this score. Traditional utilitarianism is the most obvious example of such a theory. If infinite utility is possible, and moral value is an unbounded function of utility, then infinite moral value is possible, too. We might just stipulate that utilitarianism's value function must be bounded, but what could possibly be the argument for setting the bound at one place rather than another?

This suggests that the Conceivability PEMT is an unsatisfactory substitute for the original PEMT, even if it avoids some of the original's difficulties. We must look elsewhere for a method of comparing moral value intertheoretically.

References

Broome, John. "The Most Important Thing About Climate Change," in *Public Policy: Why Ethics Matters*, ed. Jonathan Boston,

- Andrew Bradstock and David Eng. Australian National University E-Press, 2010.
- Kamm, F.M. "Terror and Collateral Damage: Are They Permissible?" *Journal of Ethics* 9 (2005), pp. 381–401.
- Lockhart, T. *Moral Uncertainty and Its Consequences*. Oxford University Press, 2000.
- Ross, Jacob. "Rejecting Ethical Deflationism," *Ethics* 116 (2006), pp. 742–768.
- Sepielli, Andrew. "What to Do When You Don't Know What to Do," *Oxford Studies in Metaethics, Volume 4*, ed. Russ Shafer-Landau. Oxford University Press, 2009.